

Exercise : geometric transform

Geometric transform in computer vision is a fundamental process aiming to compute and describe the 2D geometric transformation between 2 images. This process follows a required step of feature extractions and feature matching between the images, and has several applications such as image reconstruction, stitching and photomontage.

In the first part, the goal of this exercise is presented, then, the experimental method is explained. The obtained results are presented in a third part. After that, a discussion is conducted about the results, and a conclusion ends this report.

I - Goal

The objective of this project is to develop a program that constructs a composite image from multiple images of the same scene. This process involves aligning and blending the images to create a seamless panorama.

For this purpose, a geometric transformation of all the initial images (except one) is required. Indeed, the idea is to transform images so that they appear to have been taken from the same point. Between 2 images of the same scene, several transformations are possible : translation, euclidean, similarity, affine and projective (as presented in the image below). The more complex the transformation, the more DoF are involved, and the more parameters are required to describe the transformation.

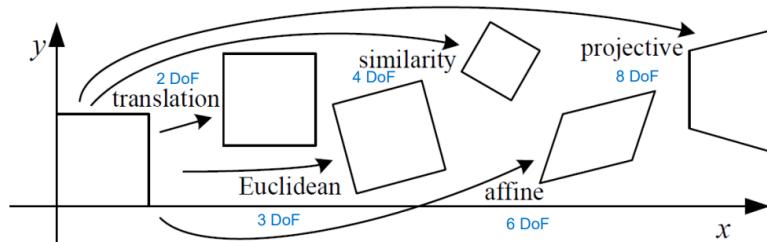


Figure 1 : basic set of 2D planar transformations

There are as many parameters characteristic of the transformation as DoF for a given transformation. It is possible to compute them from a set of matched points between both images, a set of matched points is a set of pairs of points of the same feature of the scene in both images. From this set and the characteristic Jacobian matrix of the transformation, it is possible to compute a matrix A and a vector b (as represented in the image below for the translation). Then, we obtain the parameters' vector p of the translation by solving the matrix equation : $A.p = b$.

Motion Alignment in Translation Case

$$\text{Matched Feature Points } \{(x_i, x'_i)\}_1^N = \left\{ \begin{pmatrix} x_i \\ y_i \end{pmatrix}, \begin{pmatrix} x'_i \\ y'_i \end{pmatrix} \right\}_1^N \quad \{\Delta x_i\}_1^N = \left\{ \begin{pmatrix} x'_i - x_i \\ y'_i - y_i \end{pmatrix} \right\}_1^N$$

$$\begin{aligned} \mathbf{A} &= \sum_i \mathbf{J}^T(\mathbf{x}_i) \mathbf{J}(\mathbf{x}_i) & \mathbf{A} &= \sum_{i=1}^N \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} = \begin{bmatrix} N & 0 \\ 0 & N \end{bmatrix} \\ \mathbf{b} &= \sum_i \mathbf{J}^T(\mathbf{x}_i) \Delta \mathbf{x}_i & \mathbf{b} &= \sum_{i=1}^N \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} x'_i - x_i \\ y'_i - y_i \end{bmatrix} = \begin{bmatrix} \sum_{i=1}^N (x'_i - x_i) \\ \sum_{i=1}^N (y'_i - y_i) \end{bmatrix} \\ \mathbf{Ap} &= \mathbf{b} & \begin{bmatrix} N & 0 \\ 0 & N \end{bmatrix} \begin{bmatrix} t_x \\ t_y \end{bmatrix} &= \begin{bmatrix} \sum_{i=1}^N (x'_i - x_i) \\ \sum_{i=1}^N (y'_i - y_i) \end{bmatrix} \\ && \text{2 parameters} & \end{aligned}$$

Figure 2 : computation method to obtain translation parameters

Hence, it is convenient to compute the transformation parameters between 2 images using a program script since it is mostly matrix operations. Then, the parameters can be put into a matrix, representative of the transformation. It is a 2x3 matrix for affine transformations and a 3x3 matrix for the projective transformation. However, it is possible to use a 3x3 for affine transformation for convenience by adding the third row [0 0 1]. The image below presents the transformation matrix and the Jacobian matrix of the different 2D transformations.

Jacobians of the 2D coordinate transformations

Transform	Matrix	Parameters p	Jacobian J
translation	$\begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \end{bmatrix}$	(t_x, t_y)	$\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$
Euclidean	$\begin{bmatrix} c_\theta & -s_\theta & t_x \\ s_\theta & c_\theta & t_y \end{bmatrix}$	(t_x, t_y, θ)	$\begin{bmatrix} 1 & 0 & -s_\theta x - c_\theta y \\ 0 & 1 & c_\theta x - s_\theta y \end{bmatrix}$
similarity	$\begin{bmatrix} 1+a & -b & t_x \\ b & 1+a & t_y \end{bmatrix}$	(t_x, t_y, a, b)	$\begin{bmatrix} 1 & 0 & x & -y \\ 0 & 1 & y & x \end{bmatrix}$
affine	$\begin{bmatrix} 1+a_{00} & a_{01} & t_x \\ a_{10} & 1+a_{11} & t_y \end{bmatrix}$	$(t_x, t_y, a_{00}, a_{01}, a_{10}, a_{11})$	$\begin{bmatrix} 1 & 0 & x & y & 0 & 0 \\ 0 & 1 & 0 & 0 & x & y \end{bmatrix}$
projective	$\begin{bmatrix} 1+h_{00} & h_{01} & h_{02} \\ h_{10} & 1+h_{11} & h_{12} \\ h_{20} & h_{21} & 1 \end{bmatrix}$	$(h_{00}, h_{01}, \dots, h_{21})$	(see Section 8.1.3)

re-parameterized the motions so that they are identity I for $p = \mathbf{0}$.

Figure 3 : transformation and Jacobian matrices for the different transformations

A noticeable point is that the euclidean is different from the others. Indeed, one of the transformation parameters, θ , is included in the Jacobian matrix. Hence, as for the projective transformation, it is not possible to compute the parameters with a linear method. An effective way to avoid this problem is to compute the similarity parameters, and then estimate the euclidean parameter θ with : $\theta = \tan^{-1}(a/b)$.

II - Method

In order to construct an image resulting from the appropriate blending of 2 images of the same scene, several steps are required :

- feature extraction and matching between both images
- computation of transformation matrix between the reference image and the query image
- geometric transformation of the query image
- blending the reference image and the transformed query image on a canvas

Let's detail the different steps.

a) Feature extraction and matching

This first step of the script is performed by the *matched_keypoints* function (presented in the images below), using the orb library, efficient for fastly realizing feature extractions and feature matching between images.

```
16 # find keypoints of 2 input grey scale images, perform matches between both images and return the strongest matches into an array
17 def match_keypoints(gray1, gray2, ratio, res_path):
18
19     # initiate ORB detector
20     orb = cv.ORB_create()
21
22     # kp --> Keypoints : array containing the feature points of the grayscale image
23     # des --> Descriptor : array of size [nb keypoints]x128 containing informations about feature points
24
25     # find the keypoints and compute the descriptors with ORB
26     kp1 = orb.detect(gray1, None)
27     kp2 = orb.detect(gray2, None)
28     kp1, des1 = orb.compute(gray1, kp1)
29     kp2, des2 = orb.compute(gray2, kp2)
30
31     print('nb keypoints img1 : ' + str(len(kp1)))
32
33     # draw keypoints and save the result image
34     img1_kp = gray1
35     img1_kp = cv.drawKeypoints(gray1, kp1, img1_kp)
36     cv.imwrite(res_path + 'orb_keypoints.jpg', img1_kp)
37
38     # match keypoints between both images using brute-force
39     bf = cv.BFMatcher()
40     matches = bf.knnMatch(des1, des2, k=2)
41
42     # apply ratio test to select only the strongest matches
43     good_matches = []
44     for m,n in matches:
45         if m.distance < ratio * n.distance:
46             good_matches.append([m])
47
48     # draw matches and save the result image
49     img_matches = cv.drawMatchesKnn(gray1, kp1, gray2, kp2, good_matches, None, flags=cv.DrawMatchesFlags_NOT_DRAW_SINGLE_POINTS)
50     cv.imwrite(res_path + 'orb_matches.jpg', img_matches)
51
52     # extract the 4 best matches from good_matches (those with the smallest distance)
53     if len(good_matches) < 4:
54         print("less than 4 matches between images")
55     else:
56         sorted_good_matches = sorted(good_matches, key=lambda x: x[0].distance)
57         best_4_matches = sorted_good_matches[:4]
```

```

59     # draw matches and save the result image
60     img_best_matches = cv.drawMatchesKnn(gray1, kp1, gray2, kp2, best_4_matches, None, flags=cv.DrawMatchesFlags_NOT_DRAW_SINGLE_P
61     cv.imwrite(res_path + 'orb_best_matches.jpg', img_best_matches)
62
63     # construct an array of pair of matched points
64     matched_points = []
65     for m in good_matches:
66         pt1 = kp1[m[0].queryIdx].pt
67         pt2 = kp2[m[0].trainIdx].pt
68         matched_points.append([pt1, pt2])
69
70     # construct an array of pair of the 4 best matched points
71     best_4_matched_points = []
72     for m in best_4_matches:
73         pt1 = kp1[m[0].queryIdx].pt
74         pt2 = kp2[m[0].trainIdx].pt
75         best_4_matched_points.append([pt1, pt2])
76
77     return matched_points, best_4_matched_points

```

Figure 4 : *match_keypoints* function

The *match_keypoints* function takes 2 gray images as inputs and returns a set of matched points as output. After performing the feature detection and matching using the orb library, a ratio-test is applied in order to improve the quality of the matched points in the *good_matches* variable. From this, the *matched_points* array is constructed to store only the pairs of matched points between images. Besides, another array, *best_4_matched_points* is constructed and returned by the function in order to be used by OpenCV functions. It stores only the 4 best pairs of matched points, i.e. those with the smallest “distance” in terms of feature matching.

b) Computation of transformation matrix

Once the feature matching has been performed, a set of pairs of matched points between both images is available. From this set, it is possible to compute the transformation matrix for the translation, euclidean, similarity and affine using the method presented in the *Goal* section. The following images show the different functions of the script implemented to compute the transformation parameters and matrixes.

```

120     # compute the translation parameters between 2 images using the matched keypoints and return the 3x3 transformation matrix
121     def compute_translation_matrix(matched_points):
122
123         # dimension of matrix A (biggest dim of J)
124         n = 2
125
126         # compute matrix A and vector b
127         A, b = compute_A_b(matched_points, J_translation, n)
128
129         # compute translation 2x1 vector p
130         p = solve(A, b)
131         print('translation vector :')
132         print(p)
133
134         # 3x3 transformation matrix
135         M = np.array([[1, 0, p[0][0]],
136                      [0, 1, p[1][0]],
137                      [0, 0, 1]])
138
139     return M

```

Figure 5 : *compute_translation_matrix* function

```

141 # compute the euclidean parameters between 2 images using the matched keypoints and return the 3x3 transformation matrix
142 def compute_euclidean_matrix(matched_points):
143
144     # for the estimation of the euclidean transformation, we first compute a similarity transformation since it is a non-linear problem
145
146     # dimension of matrix A (biggest dim of J)
147     ns = 4
148
149     # compute matrix A and vector b
150     As, bs = compute_A_b(matched_points, J_similarity, ns)
151
152     # compute similarity 4x1 vector p
153     ps = solve(As, bs)
154
155     # compute euclidean 3x1 vector p from the similarity vector
156     p = np.zeros([3, 1])
157     p[0][0] = ps[0][0]
158     p[1][0] = ps[1][0]
159     p[2][0] = np.arctan(ps[3][0]/ps[2][0])
160     print('euclidean vector :')
161     print(p)
162
163     # 3x3 transformation matrix
164     M = np.array([[np.cos(p[2][0]), -np.sin(p[2][0]), p[0][0]],
165                  [np.sin(p[2][0]), np.cos(p[2][0]), p[1][0]],
166                  [0, 0, 1]])
167
168     return M

```

Figure 6 : compute_euclidean_matrix function

```

170 # compute the similarity parameters between 2 images using the matched keypoints and return the 3x3 transformation matrix
171 def compute_similarity_matrix(matched_points):
172
173     # dimension of matrix A (biggest dim of J)
174     n = 4
175
176     # compute matrix A and vector b
177     A, b = compute_A_b(matched_points, J_similarity, n)
178
179     # compute similarity 4x1 vector p
180     p = solve(A, b)
181     print('similarity vector :')
182     print(p)
183
184     # 3x3 transformation matrix
185     M = np.array([[1+p[2][0], -p[3][0], p[0][0]],
186                  [p[3][0], 1+p[2][0], p[1][0]],
187                  [0, 0, 1]])
188
189     return M

```

Figure 7 : compute_similarity_matrix function

```

191 # compute the affine parameters between 2 images using the matched keypoints and return the 3x3 transformation matrix
192 def compute_affine_matrix(matched_points):
193
194     # dimension of matrix A (biggest dim of J)
195     n = 6
196
197     # compute matrix A and vector b
198     A, b = compute_A_b(matched_points, J_affine, n)
199
200     # compute affine 6x1 vector p
201     p = solve(A, b)
202     print('affine vector :')
203     print(p)
204
205     # 3x3 transformation matrix
206     M = np.array([[1+p[2][0], p[3][0], p[0][0]],
207                  [p[4][0], 1+p[5][0], p[1][0]],
208                  [0, 0, 1]])
209
210     return M

```

Figure 8 : compute_affine_matrix function

```

94 # compute the nxn matrix A and the nx1 vector b from the matched points between 2 images and the Jacobian matrix
95 def compute_A_b(matched_points, J_def, n):
96
97     A = np.zeros([n, n])
98     b = np.zeros([n, 1])
99
100    for m in matched_points:
101
102        x1 = m[0][0]
103        y1 = m[0][1]
104        x2 = m[1][0]
105        y2 = m[1][1]
106        delta = [[x2 - x1],
107                  [y2 - y1]]
108
109        # Jacobian matrix
110        J = J_def(x1, y1)
111        Jt = J.transpose()
112
113        A = A + Jt.dot(J)
114        b = b + Jt.dot(delta)
115
116    return A, b

```

Figure 9 : compute_A_b function

```

79 def J_translation(x, y):
80     J = np.array([[1, 0],
81                  [0, 1]])
82     return J
83
84 def J_similarity(x, y):
85     J = np.array([[1, 0, x, -y],
86                  [0, 1, y, x]])
87     return J
88
89 def J_affine(x, y):
90     J = np.array([[1, 0, x, y, 0, 0],
91                  [0, 1, 0, 0, x, y]])
92     return J

```

Figure 10 : implementation of Jacobian matrices

As explained in the *Goal* section, it is not possible to compute the projective transformation's matrix with a linear algorithm. However, the OpenCV function *getPerspectiveTransform* returns this matrix from the 4 pairs of matched points given as input. This solution has been implemented in the script in order to have all the transformation available.

```

225 def get_homography_matrix(matched_points):
226
227     pts1 = []
228     pts2 = []
229
230     # creation of a list of 4 points for each image
231     for m in matched_points:
232         pts1.append(m[0])
233         pts2.append(m[1])
234
235     print(pts1)
236     print(pts2)
237
238     # conversion into np array
239     pts1_np = np.array(pts1, dtype=np.float32)
240     pts2_np = np.array(pts2, dtype=np.float32)
241
242     # 3x3 transformation matrix
243     M = cv.getPerspectiveTransform(pts1_np, pts2_np)
244
245     return M

```

Figure 11 : get_homography_matrix function

c) Geometric transformation

It is possible to transform the query image using the obtained transformation matrix and the *warpPerspective* function. This is performed in the *transform* function of the script, presented in the image below. Hence, the resulting image appears to have been taken from the same point as the reference image.

```
263 # transform an image using the 3x3 transformation matrix M and save the result image
264 def transform(gray, M, res_path):
265
266     img_width = gray.shape[1]
267     img_height = gray.shape[0]
268
269     img_transformed = cv.warpPerspective(gray, M, (img_width, img_height))
270
271     print(img_transformed.shape[:2])
272
273     cv.imwrite(res_path + 'img_query_transform.jpg', img_transformed)
```

Figure 12 : transform function

d) Blending

The *transform* function only modifies the query image, but the goal is to blend the reference image and the transformed query image, in order to create a panorama. This is performed by the *blend_image* function of the script.

```
302 # stitch the reference image and the transformed query image on a single image
303 def blend_images(img_query, img_ref, M):
304
305     # compute canvas dimension
306     canvas_size, translation = calculate_canvas_size(img_query, img_ref, M)
307     tx, ty = translation
308
309     # multiplication of transformation matrix and translation matrix
310     translation_matrix = np.array([[1, 0, tx], [0, 1, ty], [0, 0, 1]])
311     M_translation = np.dot(translation_matrix, M)
312
313     # canvas creation
314     canvas = np.zeros((canvas_size[1], canvas_size[0]), dtype=np.uint8)
315
316     # put ref image on the canvas
317     canvas[ty:ty+img_ref.shape[0], tx:tx+img_ref.shape[1]] = img_ref
318
319     # transformation of the query image
320     img_query_transformed = cv.warpPerspective(img_query, M_translation, canvas_size)
321
322     # blend images on the canvas
323     canvas = cv.addWeighted(canvas, 1, img_query_transformed, 1, 0)
324
325     return canvas
```

Figure 13 : blend_images function

The *blend_images* function starts by computing the canvas dimension using the *calculate_canvas_size* function. Then, the transformation matrix is multiplied by the translation matrix computed in the *calculate_canvas_size*, the purpose of this translation is to be sure that every point's coordinates are positive. After that, a canvas of the right dimension is created and the reference image and the transformed query image are successfully added to create a stitched image.

```

275 # compute the canvas dimension required to contain the reference and the transformed query images
276 def calculate_canvas_size(img_query, img_ref, M):
277
278     # get images dimensions
279     h1, w1 = img_query.shape[:2]
280     h2, w2 = img_ref.shape[:2]
281
282     # get corners' positions for both images
283     corners_img_query = np.float32([[0, 0], [0, h1], [w1, h1], [w1, 0]]).reshape(-1, 1, 2)
284     corners_img_ref = np.float32([[0, 0], [0, h2], [w2, h2], [w2, 0]]).reshape(-1, 1, 2)
285
286     # get corners' positions of the transformed query image
287     transformed_corners_img_query = cv.perspectiveTransform(corners_img_query, M)
288
289     # combine corners of both images
290     all_corners = np.vstack((corners_img_ref, transformed_corners_img_query))
291
292     # find limits of the canvas
293     [xmin, ymin] = np.int32(all_corners.min(axis=0).ravel() - 0.5)
294     [xmax, ymax] = np.int32(all_corners.max(axis=0).ravel() + 0.5)
295     translation = [-xmin, -ymin]
296
297     # canvas dimensions
298     canvas_size = (xmax - xmin, ymax - ymin)
299
300     return canvas_size, translation

```

Figure 14 : calculate_canvas_size function

The *calculate_canvas_size* function computes the dimension of the canvas required to stitch the images. For this purpose, it calculates the transformed coordinates of the corners of the query image and computes the translation assuring the positivity of the corners.

III - Results

This section presents the results of the different parts of the script, following the order of the functions presented previously. The initial images used for the results' presentation are almost separated only by a translation of the camera, in order to make every transformation relevant between them. Hence, it is possible to compare the accuracy of the different transformations, which is performed in the *Discussion* section.



Figure 15 : initial reference image (left) and query image (right)



Figure 16 : feature detection in the query image

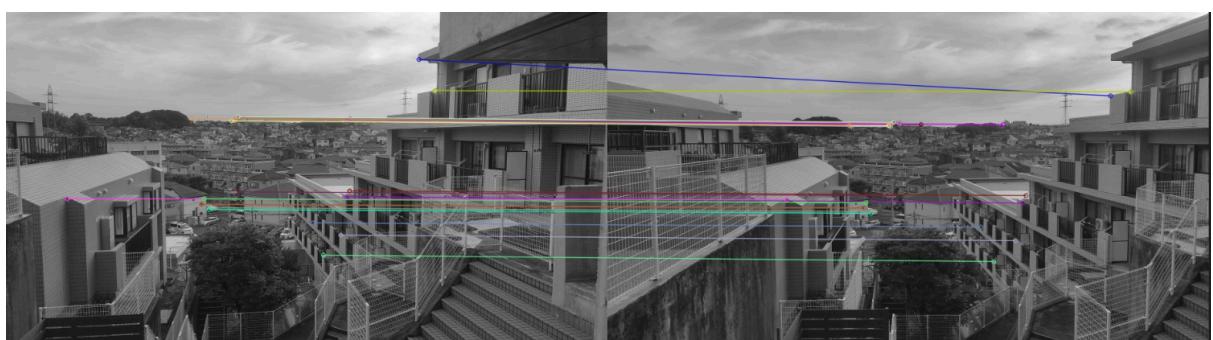


Figure 17 : feature matching between the reference image and the query image



Figure 18 : 4 best feature matching between the reference image and the query image

	Transformation matrix	Transformed query image	Stitched image
Translation	<pre>translation matrix : [[1. 0. 97.21451173] [0. 1. 8.97631469] [0. 0. 1.]]</pre>		
Similarity	<pre>similarity matrix : [[1.00910425e+00 -7.69028727e-03 9.54211708e+01] [7.69028727e-03 1.00910425e+00 3.82438215e+00] [0.00000000e+00 0.00000000e+00 1.00000000e+00]]</pre>		
Affine	<pre>affine matrix : [[1.04137586e+00 6.27287440e-02 6.62926748e+01] [2.96760428e-02 9.62063763e-01 6.13124898e+00] [0.00000000e+00 0.00000000e+00 1.00000000e+00]]</pre>		
Homography	<pre>homography matrix from func : [[5.09131526e-01 6.51604737e-02 1.52605360e+02] [-2.07580870e-01 7.93465473e-01 7.95727743e+01] [-6.37221667e-04 7.77257117e-05 1.00000000e+00]]</pre>		

Figure 19 : stitching comparison between the different transformation

It is also possible to stitch 2 images which are separated by a rotation using the script, as shown in figure 20 and 21.

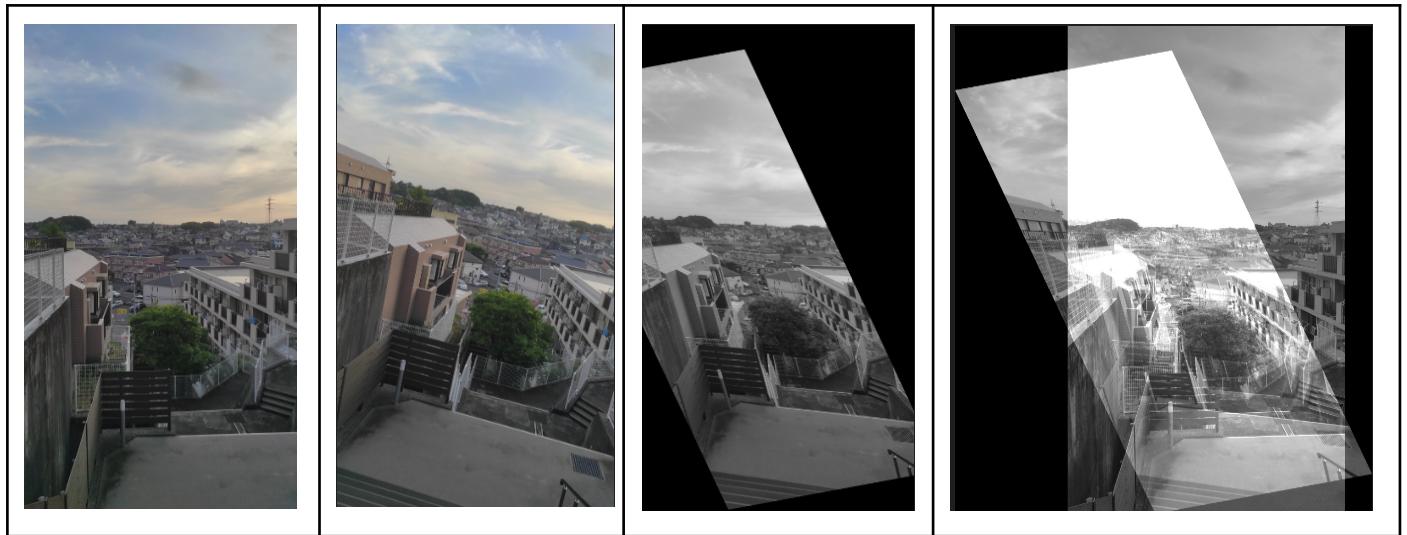


Figure 20 : stitching between 2 images using affine transformation

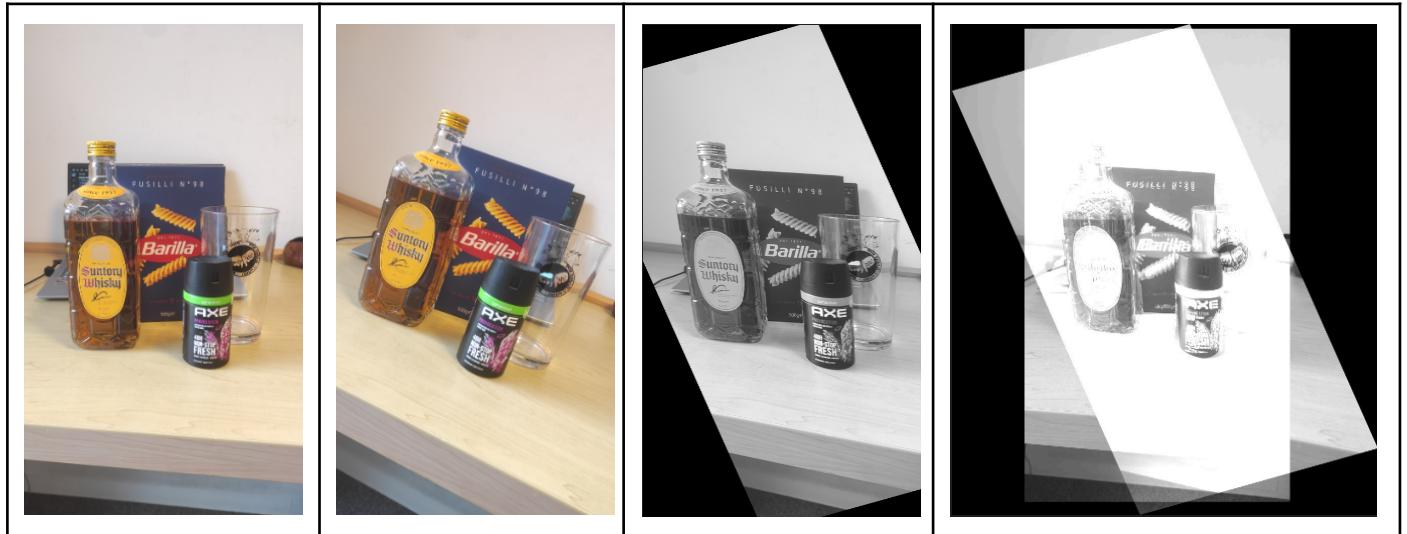


Figure 21 : stitching between 2 images using affine transformation

The different resulting images come from the developed script. Hence, it is possible to conclude that all the steps are well performed. Indeed, the final stitching works for every transformation.

An noticeable limit of the script is the homography transformation. Indeed, it is performed using the *getPerspectiveTransform* function, which takes 4 pairs of points as input and returns a transformation matrix. However, a matching error in these 4 pairs have huge consequences on the output matrix. Therefore, in several cases, the transformed query image was completely wrong because of such an error. The script used the selection of the theoretically 4 best matches (using the distances) but another method may be better.

IV - Discussion

Different geometric transformations are implemented in the script. Thus, it is possible to compare them through the exemple proposed in the *Result* section. Indeed, the query image transformation and the stitching of 2 images has been performed for the translation, the similarity transform, the affine transform and the homography. The latter's matrix has been computed through an OpenCV function while the others have been calculated using linear algebra and the Jacobian matrix.

On the figure 19, containing the different resulting images, we can see that the more complex the transformation (with the more DoF and parameters), the more accurate is the transformed query image and the stitching image. Indeed, the translation stitching looks like a draft, the different buildings are not well superposed while the homography stitching is more accurate, it is possible to clearly see that the center of the reference image has been well superposed. However, about the borders of the reference images, the superposition is not perfect for all the transformations. This is probably a consequence of the fact that the matched feature points are mainly located in the center of both images.

Hence, it is possible to conclude that more complex transformations create better stitching results. However, it becomes more complicated to compute the matrix when the complexity increases, with a particular difficulty for the homography, as this transformation required to use non-linear algebra. Moreover, the functions provided by OpenCV required a great accuracy in the input data, such as *getPerspectiveTransform* which takes only 4 pairs of matched points (as precised in the *Result* section) and is therefore non-robust to error during the feature matching process.

V - Conclusion

To conclude this exercise, I have successfully developed a python script that stitches 2 images of the same scene in order to create a panograph. For this purpose, different steps such as feature detection, feature matching, transformation matrix computation, and blending are implemented using OpenCV library and linear algebra concepts.